



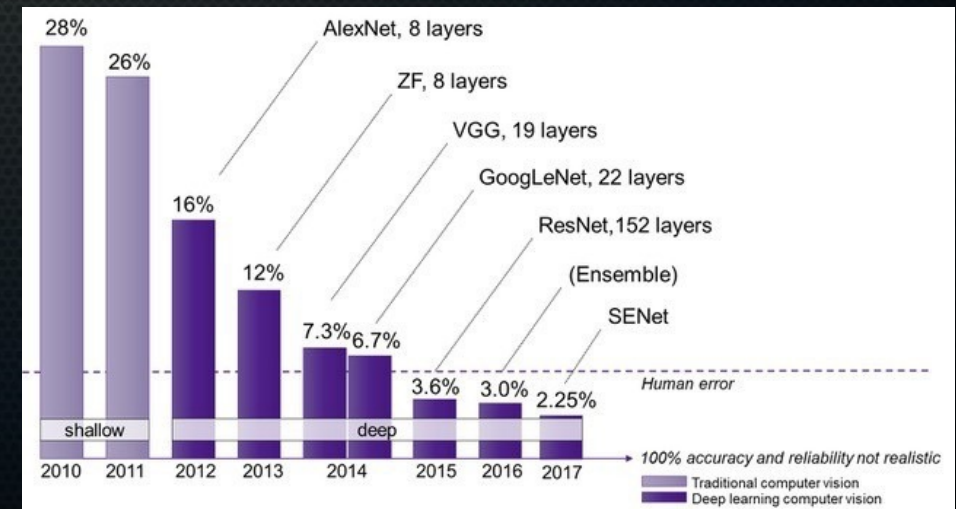

# Convolutional neural networks

- These dense neural networks are not what made the huge strides in deep learning over the last few years.



# Convolutional neural networks

- These dense neural networks are not what made the huge strides in deep learning over the last few years.
- Instead, those are deep convolutional neural networks





# Convo-what now?

---

- Let's look at an image



# Convo-what now?


- Let's look at an image
- In a dense architecture, every pixel value is connected to *every* neuron.





# Convo-what now?

- Let's look at an image
- In a dense architecture, every pixel value is connected to *every* neuron.



# Convo-what now?

- Let's look at an image
- In a dense architecture, every pixel value is connected to *every* neuron.
- This gives problems:
  - You get an *insane* amount of parameters to optimise.  $250 \times 250$  pixels \* 20 HL 1 neurons = 1,250,020 weights and biases. You can forget about any sort of findable or achievable (global) optimum.
  - There is no locality: if you want your network to know whether or not there is a dog in an image, all these parameters must be optimised so that you can recognise the dog anywhere.






# Convo-what now?

---



This is madness!



# Convo-what now?

---

- The answer: convolution. Let's look at a 1D example!



- When is the heart beating?

These slides were gladly ~~stolen~~  
borrowed (with permission) from Marc  
Pages Gallego



# Convo-what now?

---

- The answer: convolution. Let's look at a 1D example!



- When is the heart beating?

# Convo-what now?

- The answer: convolution. Let's look at a 1D example!



signal =  $x = [0 \quad 0 \quad 0 \quad 1 \quad -1 \quad 0.5 \quad 0 \quad 0 \quad 0 \quad 0 \quad 0]$

label =  $y = [0 \quad 0 \quad 0 \quad 1 \quad 1 \quad 1 \quad 0 \quad 0 \quad 0 \quad 0 \quad 0]$



# Convo-what now?


---

- Signal can be at different positions in the sequence:




# Convo-what now?

- Signal can be at different positions in the sequence:



- Dense network needs to optimise such that different weights somehow cause the network to output 1 for different positions of the signal:



label =  $y = [1 \ 1 \ 1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$




# Convo-what now?

- Signal can be at different positions in the sequence:




- Dense network needs to optimise such that different weights somehow cause the network to output 1 for different positions of the signal:




label =  $y = [1 \ 1 \ 1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$

# Convo-what now?

- Signal can be at different positions in the sequence:



- Dense network needs to optimise such that different weights somehow cause the network to output 1 for different positions of the signal:



label =  $y = [0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 1 \ 1 \ 1]$



# Convo-what now?

- Convolution:



Input

0	0	0	1	-1	0.5	0	0	0	0
---	---	---	---	----	-----	---	---	---	---

Kernel


2
---

Kernel output

0									
---	--	--	--	--	--	--	--	--	--

ReLU output

0	0	0	1	-1	0.5	0	0	0	0
---	---	---	---	----	-----	---	---	---	---



# Convo-what now?

- Convolution:




Input




\*

Kernel




Convolve (move) the  
kernel over the sequence

Kernel output



ReLU output





# Convo-what now?

- Convolution:



Input




\*


Kernel



Kernel output



ReLU output



# Convo-what now?

- Convolution:




Input




\*


Kernel



Kernel output



ReLU output





# Convo-what now?

- Convolution:



Input




\*


Kernel



Kernel output



ReLU output



# Convo-what now?

- Convolution:



Input

0	0	0	1	-1	0.5	0	0	0	0
---	---	---	---	----	-----	---	---	---	---

\*

Kernel


2

Kernel output

0	0	0	2	-2	1	0	0	0	0
---	---	---	---	----	---	---	---	---	---

ReLU output

0	0	0	2	0	1	0	0	0	0
---	---	---	---	---	---	---	---	---	---





# Convo-what now?

- Convolution:




Input



\*


Kernel



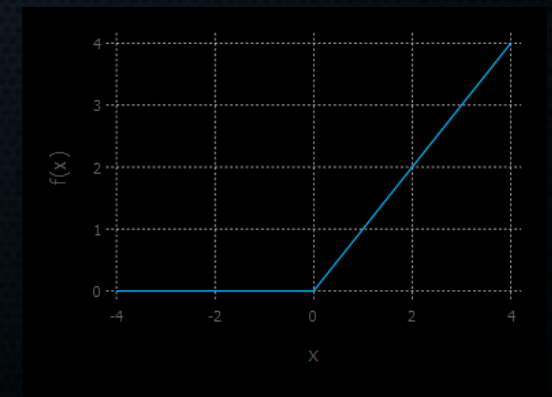
Kernel output



ReLU output



This is a kernel that detects  
positive numbers



# Convo-what now?

- Convolution:



Input




Kernel



Kernel output




ReLU output



Just like a neuron has weights, this kernel has a trainable weight (2)

This is a kernel that detects positive numbers





# Convo-what now?

- Convolution:



Input



\*

\*

\*


Kernel




A kernel can have a size  $>1$

+

Kernel output



ReLU output



# Convo-what now?

- Convolution:




Input




\* \* \*


Kernel



Kernel output



ReLU output





# Convo-what now?

- Convolution:



Input



\*

\*


\*

Kernel




+

Kernel output



ReLU output



# Convo-what now?

- Convolution:



Input




Kernel




+

Kernel output



ReLU output





# Convo-what now?


- Convolution:



Input




Kernel





+

Kernel output




Note shrinkage due to edge effects

ReLU output



# Convo-what now?

- Convolution:



# Convo-what now?

- Convolution:



Input




\*

\*


\*

Kernel




Or could compress output size further by changing *stride* (step size). Stride = 3 now.

Kernel output



ReLU output





# Convo-what now?

- Convolution:




Input




\* \* \*

Kernel





+  
}

Kernel output



ReLU output



# Convo-what now?

- Convolution:



Input

0	0	0	1	-1	0.5	0	0	0	0
---	---	---	---	----	-----	---	---	---	---



Kernel

*	*	*
0	2	1
+		


Kernel output

-	0	-	-	-1.5	-	-	0	-	-
---	---	---	---	------	---	---	---	---	---

ReLU output

-	0	-	-	0	-	-	0	-	-
---	---	---	---	---	---	---	---	---	---

Size = 3  
Stride = 3



# Convo-what now?

- Convolution:




Input



Kernel



Kernel output




ReLU output



Size = 3  
Stride = 3

This example is not useful because I randomly picked some weights for the kernel. But normally you can train these weights by backpropagation such that the network works well!





# Convo-what now?


- Convolution:



Input




Kernel



+



Kernel output



Size = 3  
Stride = 1

Note shrinkage due to edge effects

ReLU output





# Convo-what now?

Input	0	0	0	1	-1	0.5	0	0	0	0
							*	*	*	
Kernel							0	2	1	
								+		
Kernel output	-	0	1	1	-1.5	1	0	0	0	-
ReLu output	-	0	1	1	0	1	0	0	0	-

Size = 3  
Stride = 1

Not optimal, but by adding another convolution layer, you might get to something like:



# 2D convolution

- Size = 2\*2; stride = 1

0	22	0	1
12	2	3	23
3	34	26	2
0	22	86	3
4	3	1	4

0	2
2	0

68		



# 2D convolution

- Size = 2\*2; stride = 1

0	22	0	1
12	2	3	23
3	34	26	2
0	22	86	3
4	3	1	4

0	2
2	0

68	4	

# 2D convolution

- Size =  $2 \times 2$ ; stride = 1

0	22	0	1
12	2	3	23
3	34	26	2
0	22	86	3
4	3	1	4

0	2
2	0

68	4	8

# 2D convolution

- Size = 2\*2; stride = 1

0	22	0	1
12	2	3	23
3	34	26	2
0	22	86	3
4	3	1	4

0	2
2	0

68	4	8
10		

Etc.



# Another type of convolution: max pooling

- Size =  $2 \times 2$ ; stride = 1; just take the maximum value in the kernel area

0	22	0	1
12	2	3	23
3	34	26	2
0	22	86	3
4	3	1	4

22	22	23
34	26	26
34	86	86
22	86	86

# Another type of convolution: pooling/averaging

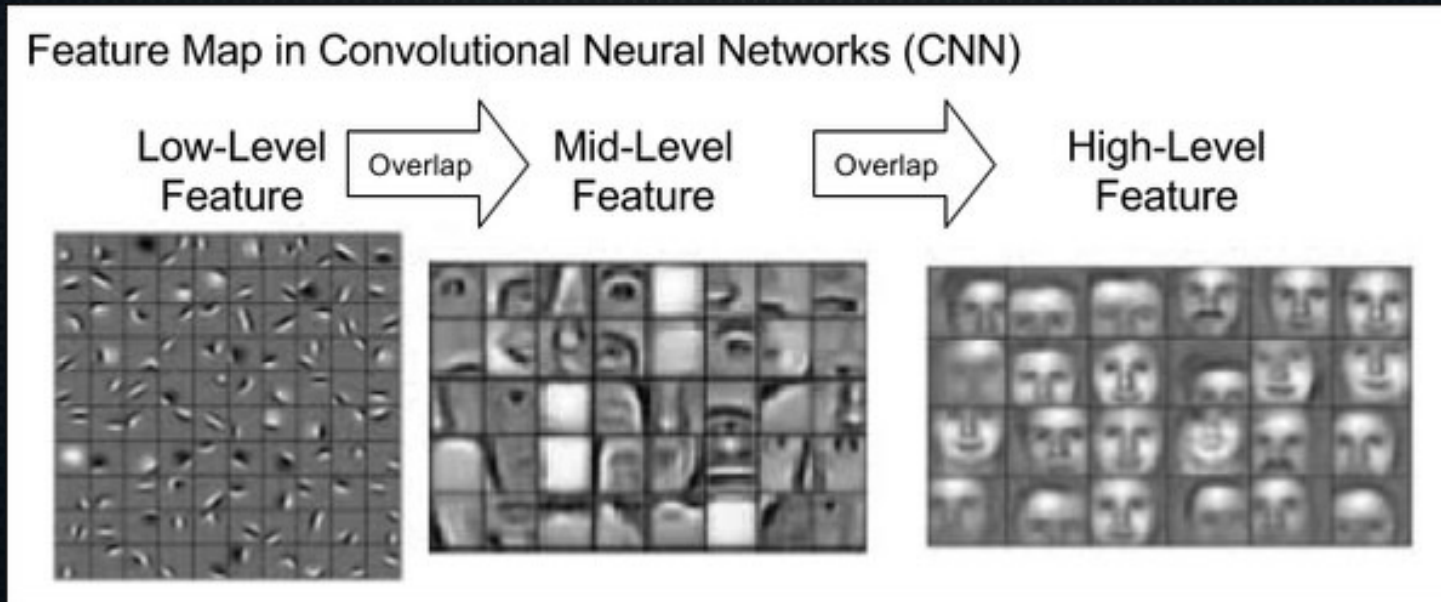
---






# Use in face detection

- Since kernels so few parameters: can use *many* of them per layer → each becomes sensitive to different image features

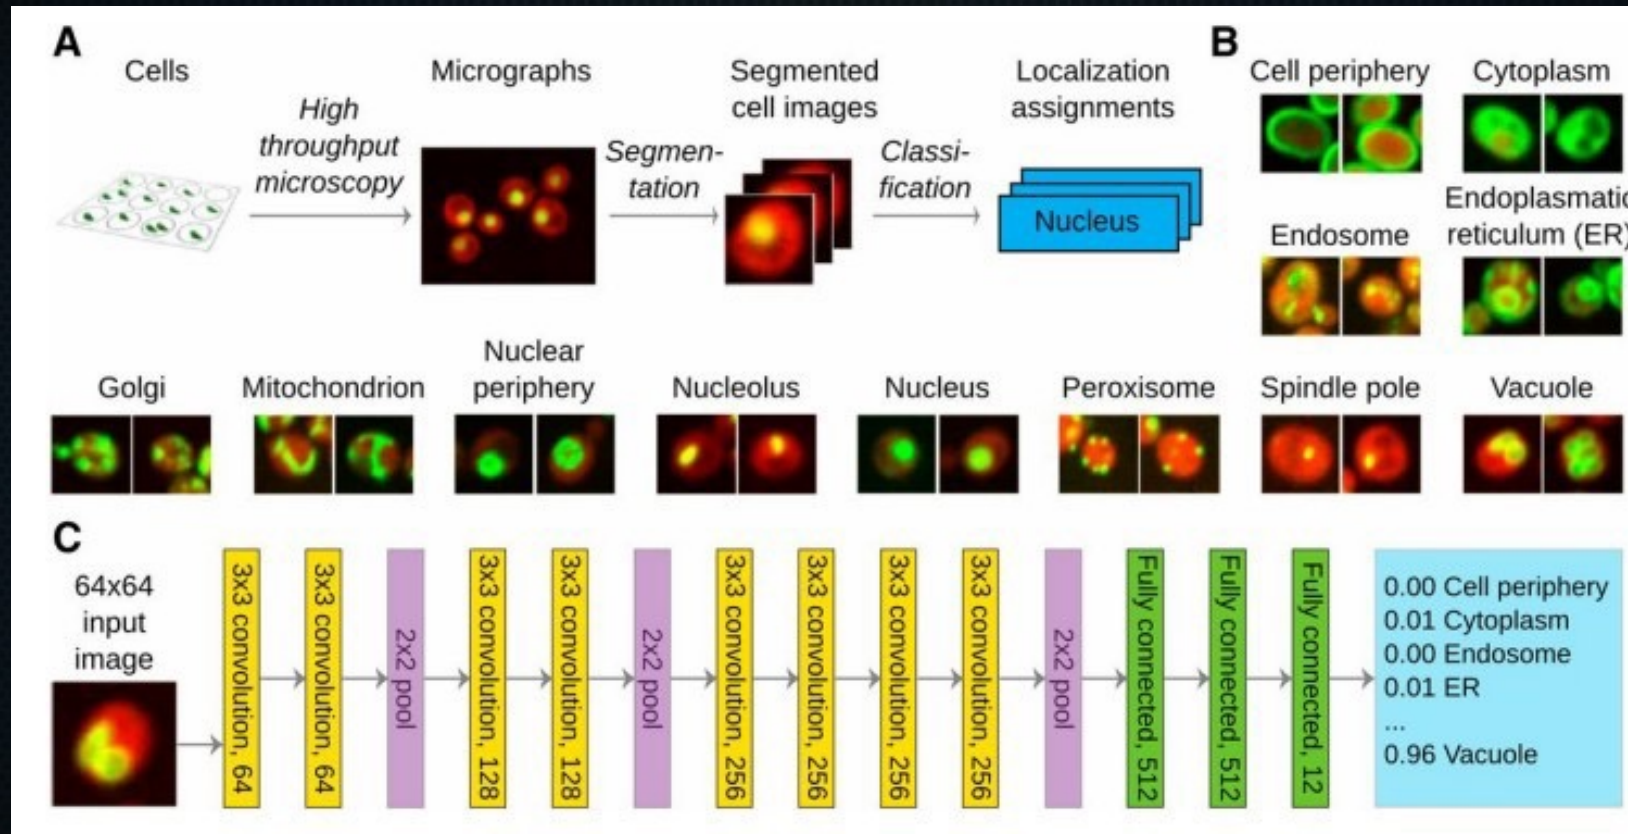




# Example AlexNet (2012)



# Biological use



Pärnamaa, T., & Parts, L. (2017). Accurate classification of protein subcellular localization from high-throughput microscopy images using deep learning. *G3: Genes, Genomes, Genetics*, 7(5), 1385-1392.



# Biological use

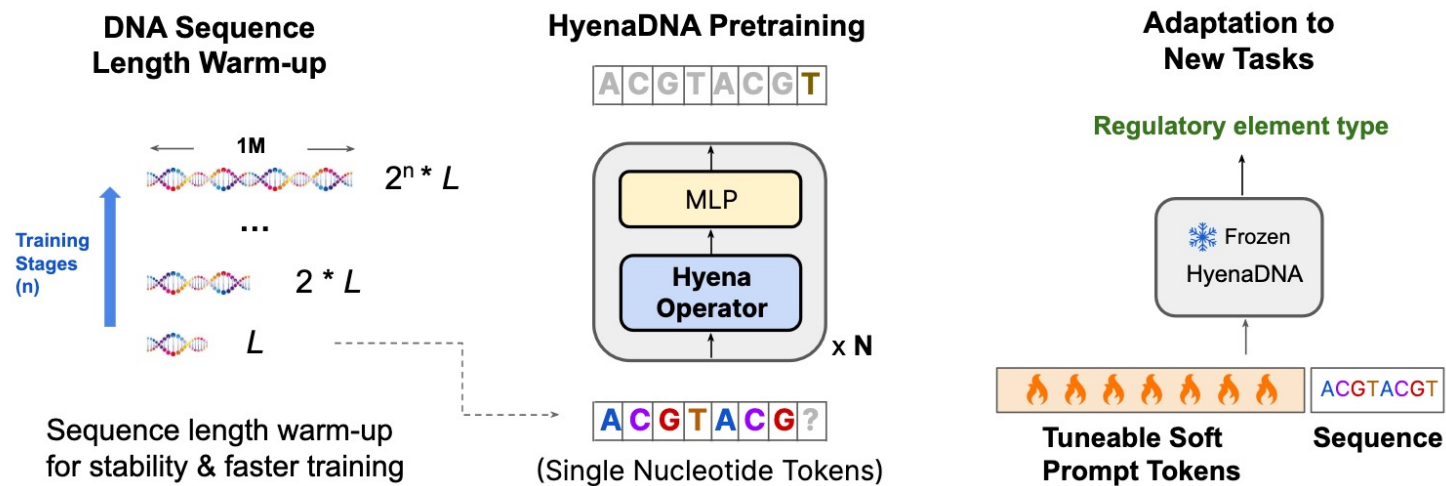



Figure 1.1: HyenaDNA recipe for long-range foundation models in genomics. The HyenaDNA architecture is a simple stack of Hyena operators (Poli et al., 2023) trained using next token prediction. (See Fig. 1.3 for block diagram of architecture). We introduce a new sequence length scheduling technique to stabilize training, and provide a method to leverage the longer context length to adapt to novel tasks without standard fine-tuning by filling the context window with learnable soft prompt tokens.



# There's a lot more

- Batch normalisation
- Vanishing gradient problem
- Dropout
- Recurrent neural nets



# Implementation

---

- We are not going to implement convolutional neural networks ourselves: implementing backpropagation properly on a simple dense network is already taxing enough.
- Still, doing that should give you a solid basis for understanding convolutional neural networks, and we'll introduce the Keras library for building (convolutional) neural networks next Monday.

# Afternoon practical

---

- Implement backpropagation yourself
- Train a dense neural network on the MNIST dataset